

PRAHARI

AI-Powered Digital Public Safety Command Centre
Predict. Protect. Prosecute.

ET AI Hackathon 2.0 — Phase 2: Build Sprint — Prototype Submission

Problem Statement #6: AI for Digital Public Safety —
Defeating Counterfeiting, Fraud & Digital Arrest Scams

1. The Problem: An Industrial-Scale Assault on Digital Trust

India is facing two converging threats to public trust in its financial and digital systems. The first is the **digital arrest scam**: organised, cross-border fraud compounds impersonate CBI, police, RBI and Enforcement Directorate officials over video calls — using spoofed numbers, AI-generated voices, deepfake video and forged government credentials — to hold victims in psychological custody for hours and coerce them into transferring their savings to 'verification accounts'. In 2024 alone, the Indian Cyber Crime Coordination Centre (I4C) recorded over **92,000 digital arrest complaints with losses exceeding Rs 1,935 crore**, with senior citizens disproportionately targeted.

The second is **Fake Indian Currency Notes (FICN)**: high-quality counterfeit Rs 500 notes routinely bypass manual detection at local cash counters, small businesses and even bank branches, entering circulation through known border corridors. Detection today depends on touch, instinct and expensive fixed machines that never reach the last-mile cash economy.

The deeper failure is structural: **every response today is reactive and siloed**. An FIR is filed after the money is gone; a counterfeit is discovered after it has changed hands; and the coordinated campaign behind hundreds of 'unrelated' cases across states remains invisible because no system connects victim reports, call infrastructure, device fingerprints and mule accounts into one picture.

2. The Solution: PRAHARI

PRAHARI (Hindi: 'sentinel') is a unified AI intelligence platform for law-enforcement agencies, financial institutions and citizens that shifts digital public safety from **reactive case investigation to predictive threat neutralisation**. One command centre, three AI engines:

Module	Who it protects	What it does
Scam Shield NLP scam classifier	Citizens, telecom providers	Classifies live call transcripts against digital-arrest playbooks in under two seconds and alerts the victim before the transfer occurs , with automated telecom spoof-block requests and MHA/I4C incident reports.
FICN Scanner Computer vision	Bank tellers, merchants, field officers	Turns any smartphone camera, counting machine or PoS terminal into a currency authenticator — verifying seven RBI security features and cross-checking serial numbers against the national seizure registry.
Network Intel Graph AI	Police cyber cells, I4C, banks	Fuses victim reports, VoIP signatures, device fingerprints and mule-account linkages into one graph; community detection exposes the coordinated campaign across jurisdictions and exports a SHA-256-sealed, court-ready evidence package.

These engines feed a **National Command Centre**: a live geospatial threat map of fraud hotspots, scam-call origin zones and FICN entry corridors, with a real-time fused intelligence feed for command staff.

3. System Architecture

The prototype is a client-side React application; the architecture below shows the production data flow the prototype models end-to-end.

INGEST	AI ENGINES	ACTION
Telco call metadata & live transcripts (SIP/STT)	Scam Shield — weighted multi-category NLP signal fusion (authority impersonation, victim isolation, urgency/fear, financial extraction, identity harvesting); LLM-augmentable	Victim push-alert · telecom spoof-block · auto-generated MHA/I4C report
Smartphone / counting-machine / PoS image capture	FICN Scanner — CV verification of security thread, watermark & electrotpe, microlettering, intaglio relief, see-through register, latent image, serial-pattern validation	Instant teller verdict · seizure report pre-filled for NCRB FICN registry · batch/corridor linkage
I4C victim reports · CDR metadata · bank device fingerprints · account linkages	Network Intel — entity resolution, Louvain community detection, temporal correlation across 38+ entity graph	Campaign confirmation · mule-account freeze drafts · inter-state alerts · SHA-256-sealed evidence package

Technology stack

- **Frontend / command centre:** React 18 + Vite, Leaflet geospatial mapping, canvas-based force-directed graph rendering (d3-force).
- **Scam classification:** custom weighted signal-fusion NLP engine over five behavioural categories derived from I4C-documented scam playbooks; designed to be augmented with a fine-tuned transformer / LLM scorer in production.
- **Counterfeit detection:** region-of-interest feature verification pipeline across the RBI security-feature set, with confidence scoring per feature and serial-number registry cross-match.
- **Graph intelligence:** typed entity graph (kingpin, operator cells, VoIP blocks, SIM boxes, device fingerprints, mule accounts, victims) with community detection and evidence-package export sealed by SHA-256 (Web Crypto).
- **Integrity & auditability:** every exported intelligence package carries an integrity seal and chain-of-custody metadata for legal admissibility.

4. Prototype Walkthrough (What the Demo Shows)

- **Command Centre:** national threat map with live hotspots (Jamtara, Nuh, Malda corridor, metro clusters) and a fused real-time intelligence feed.
- **Scam Shield:** a real digital-arrest script pattern is analysed live — risk score 97/100, verdict 'DIGITAL ARREST SCAM DETECTED', every manipulative phrase highlighted by signal category; a genuine bank service call scores near zero, demonstrating the low false-positive behaviour required

of citizen-facing tools.

- **FICN Scanner:** a suspect Rs 500 specimen fails 6 of 7 security features (14.6% authenticity); the serial number is matched to a known seized batch from the Malda corridor and a seizure report is pre-filled.

- **Network Intel:** twelve seemingly unrelated victim reports across seven states are correlated into a single confirmed campaign ('OPX-2231 — Operation Parcel Trap', confidence 94%) 31 hours after the first report, with Rs 4.2 crore of flow mapped through six mule accounts and a court-ready evidence package exported in one click.

5. Mapping to the Evaluation Focus

Evaluation criterion	How PRAHARI addresses it
Digital arrest scam detection precision & recall	Weighted fusion across five independent signal categories — no single keyword can trigger a critical verdict; benign calls demonstrably score low. Categories map to documented I4C playbook stages, giving explainable verdicts.
Counterfeit detection across denominations & print quality	Per-feature verification with individual confidence scores generalises across the RBI security-feature set; serial-pattern validation catches even high-print-quality counterfeits that reuse or increment seized series.
Fraud network detection lead time	Graph correlation confirms a coordinated campaign at T+31h after the first victim report — while the campaign is still recruiting victims — versus months of siloed, per-FIR investigation.
Very low false positives for citizen-facing tools	Saturating category scores, verdict tiers and explainable per-signal breakdowns; the demo explicitly includes a genuine bank call as a negative control.
Auditability / legal admissibility	Evidence packages are SHA-256 integrity-sealed, timestamped, and carry chain-of-custody metadata (source systems, lawful-intercept references).

6. Impact & Why This Matters

- **Prevention, not post-mortem:** the single highest-leverage moment in a digital arrest scam is the window before the first transfer. Scam Shield is built for exactly that window.
- **Last-mile currency trust:** FICN detection moves from expensive fixed infrastructure to every smartphone — reaching kirana stores, petrol pumps and rural bank correspondents.
- **From cases to campaigns:** Network Intel gives police what fraud compounds fear most — cross-jurisdiction visibility, fast enough to freeze mule accounts while money is still in the pipeline.
- **Citizen reach:** the roadmap extends Scam Shield to WhatsApp/IVR in 12 regional languages, guided reporting to the 1930 helpline and NCRB portals.

7. Production Roadmap

- **Phase 1 (0–3 months):** pilot Scam Shield with one telecom circle using on-device speech-to-text; fine-tune transformer classifier on I4C complaint corpus; FICN model training on RBI/NCRB seizure imagery.
- **Phase 2 (3–9 months):** bank integrations (counting machines, teller apps); I4C complaint-pipeline ingestion for Network Intel; evidence-package format alignment with BNSS digital-evidence rules.

- **Phase 3 (9–18 months):** national command-centre deployment, deepfake/AI-voice detection module, 12-language Citizen Shield on WhatsApp and IVR.

Note: all statistics displayed inside the prototype UI, and all case entities, are synthetic demonstration data modelled on publicly reported I4C/NCRB patterns. The currency specimen is a stylised SVG render, not an image of real currency. Public figures cited in Section 1 are from press reporting of I4C data.

PRAHARI — built for ET AI Hackathon 2.0, Phase 2 Build Sprint. Repository and 3–4 minute demo video accompany this document.